

# CodeReef: an open platform for portable MLOps, reusable automation actions and reproducible benchmarking

Grigori Fursin, Hervé Guillou, Nicolas Essayan

CodeReef.ai

January, 2020

## Abstract

We present CodeReef - an open platform to share all the components necessary to enable cross-platform MLOps (MLSysOps), i.e. automating the deployment of ML models across diverse systems in the most efficient way. We also introduce the CodeReef solution - a way to package and share models as non-virtualized, portable, customizable and reproducible archive files. Such ML packages include JSON meta description of models with all dependencies, Python APIs, CLI actions and portable workflows necessary to automatically build, benchmark, test and customize models across diverse platforms, AI frameworks, libraries, compilers and datasets.

We demonstrate several CodeReef solutions to automatically build, run and measure object detection based on SSD-Mobilenets, TensorFlow and COCO dataset from the latest MLPerf inference benchmark across a wide range of platforms from Raspberry Pi, Android phones and IoT devices to data centers. Our long-term goal is to help researchers share their new techniques as production-ready packages along with research papers to participate in collaborative and reproducible benchmarking, compare the different ML/software/hardware stacks and select the most efficient ones on a Pareto frontier using online CodeReef dashboards.

**Keywords:** *portable MLOps, MLSysOps, portable workflows, reproducibility, reusability, automation, benchmarking, optimization, collective knowledge format*

## 1 Introduction

While helping different organizations to deploy machine learning models in production during past 10 years we have noticed that finding the relevant code and training models is only a tip of the MLOps iceberg. The major challenge afterwards is to figure out how to integrate such models with complex production systems and run them in the most reliable and efficient way across rapidly evolving software, heterogeneous hardware and legacy platforms with continuously changing interfaces and data formats while balancing multiple characteristics including speed, latency, accuracy, memory size, power consumption, reliability and costs.

Different tools were introduced in the past few years to cope with this very tedious, ad-hoc and error-prone process:

- ML workflow frameworks such as MLFlow [23], Kedro [21] and Amazon SageMaker [1] help to abstract and automate high-level ML operations. However unless used inside AWS or DataBricks cloud they still have a limited support for the complex system integration and optimization particularly when tar-

geting embedded devices and IoT - the last mile of MLOps.

- ML benchmarking initiatives such as MLPerf [22], MLModelScope [19] and Deep500 [15] attempt to standardize benchmarking and co-design of models and systems. However integration with complex systems and adaptation to continuously changing user environments and data is out of their scope.
- Package managers such as Spack [17] and Easy-Build [18] are very useful to rebuild the whole environment with fixed software versions. However adaptation to existing environment, native cross-compilation and support for non-software packages (models, data sets) is still in progress.
- Docker, Kubernetes and other container-based technology is very useful to prepare and share stable software releases. However, it hides all the software chaos rather than solving it, has some performance overheads and requires enormous amount of space, have a very poor support for embedded devices and do not help to integrate models with native environment and user data.



on a user machine or install/cross-compile the missing ones while supporting different operating systems (Linux, Windows, MacOS, Android) and hardware (Nvidia, Arm, Intel, AMD ...).

Such approach allows researchers to create and share flexible APIs with JSON input/output for different AI/ML frameworks, libraries, compilers, models and datasets, connect them together into unified workflows instead of hardwired scripts, and make them portable [13] using the automatic software detection plugins [12] and meta-packages [11].

We implemented CodeReef platform with an open-source CodeReef client [7] to solve two major drawbacks of CK: it is now possible to publish CK components with different versions on a platform, and to download stable versions with all dependencies for a given workflow. We also provided an API in the CodeReef client to init, run and validate the so-called CodeReef solutions with ML models based on JSON manifest describing all CK dependencies and installation/compilation recipes for different target platforms.

We believe that combining CodeReef and CK can help users to implement and share portable ML workflows assembled from stable and versioned CK components, keep track of all the information flow within such workflows, expose configuration and optimization parameters from different tools and models and change them via JSON input files, combine shared and user code and data, monitor system behavior, retarget ML models with all the necessary software to different platforms from IoT to cloud, use them inside containers, integrate them with legacy systems and reproduce all exposed characteristics - all the pillars of practical MLOps.

### 3 CodeReef demo: automating, sharing and reproducing MLPerf inference benchmarks

A prototype of our platform is now available at [CodeReef.ai/portal](https://codereef.ai/portal) and we started working with the community and different organizations to share ML models from research papers and standard benchmarks as portable CodeReef solutions, collaboratively benchmark them across diverse platforms and reproduce performance numbers [14].

For example, as a part of the MLPerf benchmarking consortium [2], we want to automate the manual and tedious process of submitting MLPerf inference benchmark results and share MLPerf models as portable CodeReef solutions. Since dividiti, one of the MLPerf submitters, already used open-source CK workflows to submit MLPerf inference results [3], we demonstrate how to convert them into stable CodeReef solutions with portable ML models, share them with the community, connect them with the CodeReef dashboard and crowdsourcing

benchmarking across diverse platforms provided by volunteers similar to SETI@home at [CodeReef.ai/demo](https://codereef.ai/demo).

This live and interactive demo shows how to use the CodeReef solution with a unified API to automatically build, run and validate object detection based on SSD-MobileNets, TensorFlow and COCO dataset on Raspberry Pi, Android phones, laptops, desktops and data centers. When preparing this solution we had to manually create a JSON file (we plan to provide a GUI to simplify this process) describing all the dependencies on CK components and APIs to automate the following tasks:

- prepare a Python virtual environment (can be skipped for the native installation),
- download and install the Coco dataset (50 or 5000 images),
- detect C++ compilers or Python interpreters needed for object detection,
- install Tensorflow framework with a specified version for a given target machine,
- download and install the SSD-MobileNet model compatible with selected Tensorflow,
- manage installation of all other dependencies and libraries,
- compile object detection for a given machine and prepare pre/post-processing scripts.

We published this solution at the CodeReef platform using the open-source CodeReef client [7] to let the users download, initialize, run it and participate in crowd-benchmarking using their machines as follows:

1. install CodeReef client from PyPi using:

```
pip install codereef
```

2. download and install the solution on a given machine (example for Linux):

```
cr init demo-obj-detection-coco-tf-cpu-benchmark-linux-portable-workflows
```

3. run the solution on a given machine:

```
cr benchmark demo-obj-detection-coco-tf-cpu-benchmark-linux-portable-workflows
```

The users can then see all their measurements (speed, latency, accuracy) and compare them against the official MLPerf results using the CodeReef dashboard for this solution at [codereef.ai/portal/c/cr-result/sota-mlperf-object-detection-v0.5-crowd-benchmarking](https://codereef.ai/portal/c/cr-result/sota-mlperf-object-detection-v0.5-crowd-benchmarking).

After validating this solution on a given platform, the users can retarget it to other devices and operating systems such as MacOS, Windows with Docker, Android phones, servers with CUDA-enabled GPUs and so on.

We also demonstrate how the unified API of CodeReef solutions can help to integrate ML models with legacy systems and applications by integrating the above workflow with the live object detection from the web cam available at [CodeReef.ai/portal/c/cr-solution/demo-obj-detection-coco-tf-cpu-webcam-linux-azure](https://codereef.ai/portal/c/cr-solution/demo-obj-detection-coco-tf-cpu-webcam-linux-azure).

## 4 Conclusions

We presented CodeReef - an open platform to share ML models and related artifacts from research papers and benchmarks as non-virtualized, portable, customizable and reusable open-source packages while automating their benchmarking and deployment in production across diverse platforms and datasets. Since we are only at the beginning of this long-term community project, we collaborate with the community, companies, ACM, MLPerf and several systems and machine learning conferences to improve the CodeReef platform and support their reproducibility initiatives, automate ML benchmarks and enable portable MLOps based on real-world use cases.

## References

- [1] Amazon SageMaker: fully managed service to build, train, and deploy machine learning models quickly. <https://aws.amazon.com/sagemaker>.
- [2] MLPerf: a broad ML benchmark suite for measuring performance of ML software frameworks, ML hardware accelerators, and ML cloud platforms. <https://mlperf.org>.
- [3] MLperf inference benchmark results. <https://mlperf.org/inference-results>.
- [4] Artifact Evaluation: reproducing experimental results from published research papers at systems and machine learning conferences (methodology and automation tools). <https://cTuning.org/ae>, 2015–current.
- [5] Collective Knowledge use cases. <https://cKnowledge.org/partners>, 2015–present.
- [6] Collective Knowledge(CK): an open-source and cross-platform framework to decompose complex software projects into portable, customizable and reusable components and workflows with Python APIs, JSON input/output files and an integrated package manager. <https://github.com/ctuning/ck>, 2015–present.
- [7] Cross-platform CodeReef client: a command line interface and a python api to communicate with the codereef platform. <https://github.com/codereef-ai/client>, 2015–present.
- [8] Research papers with portable CK workflows. <https://codereef.ai/portal/search?q=reproduced-papersANDportable-workflow-ck>, 2015–present.
- [9] Shared and versioned CK modules. <https://codereef.ai/portal/c/module>, 2015–present.
- [10] Shared CK API for R&D automation actions and best practices. <https://codereef.ai/portal/c/cr-action>, 2015–present.
- [11] Shared CK meta packages with automated installation across diverse platforms. <https://codereef.ai/portal/c/package>, 2015–present.
- [12] Shared CK plugins to detect software (models, frameworks, data sets, scripts) on a user machine. <https://codereef.ai/portal/c/soft>, 2015–present.
- [13] Shared portable CK workflows. <https://codereef.ai/portal/c/program>, 2015–present.
- [14] Portable CodeReef solutions: sharing ML models as portable, customizable and reusable packs with a standardized API and JSON input/output. <https://CodeReef.ai/portal/c/cr-solution>, 2019–current.
- [15] T. Ben-Nun, M. Besta, S. Huber, A. N. Ziogas, D. Peter, and T. Hoefer. A Modular Benchmarking Infrastructure for High-Performance and Reproducible Deep Learning. IEEE, May 2019. The 33rd IEEE International Parallel & Distributed Processing Symposium (IPDPS'19).
- [16] Grigori Fursin, Anton Lokhmotov, and Ed Plowman. Collective Knowledge: towards R&D sustainability. In *Proceedings of the Conference on Design, Automation and Test in Europe (DATE'16)*, March 2016.
- [17] T. Gamblin, M. LeGendre, M. R. Collette, G. L. Lee, A. Moody, B. R. de Supinski, and S. Futral. The spack package manager: bringing order to hpc software chaos. In *SC '15: Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–12, Nov 2015.
- [18] Kenneth Hoste, Jens Timmerman, Andy Georges, and Stijn Weirtdt. Easybuild: Building software with ease. pages 572–582, 11 2012.
- [19] Cheng Li, Abdul Dakkak, Jinjun Xiong, Wei Wei, Lingjie Xu, and Wen mei Hwu. Across-stack profiling and characterization of machine learning models on gpus, 2019.
- [20] General Motors. Collaboratively Benchmarking and Optimizing Deep Learning Implementations. <https://youtu.be/1ldgVZ64hEI>, 2017.

- [21] QuantumBlack. Kedro: the open source library for production-ready machine learning code. <https://github.com/quantumblacklabs/kedro>, 2019.
- [22] Vijay Janapa Reddi, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, Carole-Jean Wu, Brian Anderson, Maximilien Breughe, Mark Charlebois, William Chou, Ramesh Chukka, Cody Coleman, Sam Davis, Pan Deng, Greg Diamos, Jared Duke, Dave Fick, J. Scott Gardner, Itay Hubara, Sachin Idgunji, Thomas B. Jablin, Jeff Jiao, Tom St. John, Pankaj Kanwar, David Lee, Jeffery Liao, Anton Lokhmotov, Francisco Massa, Peng Meng, Paulius Micikevicius, Colin Osborne, Genady Pekhimenko, Arun Tejusve Raghunath Rajan, Dilip Sequeira, Ashish Sirasao, Fei Sun, Hanlin Tang, Michael Thomson, Frank Wei, Ephrem Wu, Lingjie Xu, Koichi Yamada, Bing Yu, George Yuan, Aaron Zhong, Peizhao Zhang, and Yuchen Zhou. Mlperf inference benchmark, 2019.
- [23] Matei Zaharia, Andrew Chen, Aaron Davidson, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, Siddharth Murching, Tomas Nykodym, Paul Ogilvie, Mani Parkhe, Fen Xie, and Corey Zumar. Accelerating the machine learning lifecycle with mlflow. *IEEE Data Eng. Bull.*, 41:39–45, 2018.